

Practical: AWS S3 Multi-Part Upload (Manual)

CLI Objective: To manually execute the S3 Multi-Part Upload process using the AWS Command Line Interface (CLI), breaking a large file into parts, uploading them independently, and then completing the upload for reassembly by S3.

This method is recommended for object sizes greater than **100 MB**.

Phase 1: Prepare the File (Manual Split)

Unlike using an AWS SDK, when using the CLI manually, **Pranav** must first split the file into smaller chunks using a system command like `split`.

1. **Generate a Large File:** Ensure a large file (e.g., 1.5 GB) exists on your local machine (e.g., `order-items-new-data-set.csv`).
2. **Split the File:** Use the `split` command to break the file into parts of a specified size (e.g., 100 MB). This is a Linux/Unix command, not an S3 command.

Bash

```
# Example command to split a file into 100MB parts, naming the output files x00, x01, etc.
```

```
$ split -b 100m order-items-new-data-set.csv x  
# This creates output files like x00, x01, x02...
```

Phase 2: Initiate the Upload and Upload Parts

This phase initiates the process and uploads the individual parts in parallel.

1. **Initiate Multi-Part Upload:** Run the `create-multipart-upload` command, specifying the S3 bucket and the key (final file name).
 - **Result:** AWS S3 returns a unique **Upload ID**, which is a crucial identifier for all remaining steps.

Bash

```
$ aws s3api create-multipart-upload --bucket pract-5-pranav --key dataset.csv  
# Output includes the unique "UploadId"
```

2. **Upload Individual Parts:** For each split file part (e.g., `x00`, `x01`), run the `upload-part` command, specifying the unique metadata for that part.
 - **CRITICAL:** This command must include the **Part Number**, the local file name in the `--body` parameter, and the **Upload ID**.

```

Bash
# Upload Part 1 (x00)
$ aws s3api upload-part --bucket pract-5-pranav --key dataset.csv --part-number 1 --body xac
--upload-id
WnUruKVhwZNRo.TTNuK6MMOnM0qUD.5Exgrosztlh5ic9zYLTUHoXNCR6omk4c1amNnxZ92
zPcZlcv0ahvHDrnZJU7CCHruo0CQXS2yFn2JqmXOTOFIzJqX6bHQhIQg

# Upload Part 2 (x01)
$ aws s3api upload-part
--bucket <your-bucket-name>
--key order-items-new-data-set.csv
--part-number 2
--body x01
--upload-id <Your-Upload-ID>
# Repeat for all 16 parts

```

3.

- **Note:** If any part fails, only that part needs to be retried, not the entire upload.
Each successful upload returns an **ETag**.

Phase 3: Complete the Upload (Assemble the File)

After all parts are uploaded, you must instruct S3 to assemble them into the final object.

List Parts to Collect ETags: Run the [list-parts](#) command to retrieve the list of all successfully uploaded parts, including their required ETag values. The output must be formatted to include only the Part Number and ETag in a JSON format.

Bash

This command uses the required query option to format the output as needed.

```
$ aws s3api list-parts --bucket pract-5-pranav --key dataset.csv --upload-id
WnUruKVhwZNRo.TTNuK6MMOnM0qUD.5Exgrosztlh5ic9zYLTUHoXNCR6omk4c1amNnxZ92
zPcZlcv0ahvHDrnZJU7CCHruo0CQXS2yFn2JqmXOTOFIzJqX6bHQhIQg --query '{Parts:
Parts[].{PartNumber:PartNumber,ETag:ETag}}' > output.json
```

The output is saved to output.json

Complete the Upload: Use the [complete-multipart-upload](#) command, referencing the [output.json](#) file to tell S3 which parts to combine.

Replace output.json with

```
{
  "Parts": [
    {
      "PartNumber": 1,
      "ETag": "WnUruKVhwZNRo.TTNuK6MMOnM0qUD.5Exgrosztlh5ic9zYLTUHoXNCR6omk4c1amNnxZ92zPcZlcv0ahvHDrnZJU7CCHruo0CQXS2yFn2JqmXOTOFIzJqX6bHQhIQg"
    }
  ]
}
```

```

    "PartNumber": 1,
    "ETag": "\"etag1\""
},
{
    "PartNumber": 2,
    "ETag": "\"etag2\""
},
{
    "PartNumber": 3,
    "ETag": "\"etag3\""
}
]
}
}

```

Bash

```
$ aws s3api complete-multipart-upload --bucket pract-5-pranav --key dataset.csv --upload-id
WnUruKVhwZNRo.TTNuK6MMOnM0qUD.5Exgroszlh5ic9zYLTUHoXNCR6omk4c1amNnxZ92
zPcZlcv0ahvHDrnZJU7CCHruo0CQXS2yFn2JqmXOTOFIzJqX6bHQhIQg --multipart-upload
file://output.json
```

1.

- **Validation:** The upload is now complete, and the final 1.5 GB file appears in the S3 bucket.

Phase 4: Cleanup (Abort Non-Completed Uploads)

If an upload fails or is canceled before completion, you must explicitly abort it to avoid storage charges for orphaned parts.

Abort Command: Use the [abort-multipart-upload](#) command with the Upload ID of the non-completed upload.

Bash

```
$ aws s3api abort-multipart-upload \
--bucket <your-bucket-name> \
--key <Non-Completed-File-Key> \
--upload-id <Non-Completed-Upload-ID>
```

1.

- This action removes all parts associated with that Upload ID.